Conditional variance: samples behind the comparison

50,000 optimizer updates | all five held-out seeds | original KSIVI loss

Shared conditional

Equal learning rates

$\eta_v = 10^{-3}$

Mean SW2: 0.870 +/- 0.524

Shared conditional

Slow variance head

$\eta_v = 10^{-5}$

Mean SW2: 0.129 +/- 0.044

Independent linear

Variance map

$\eta_v = 10^{-3}$

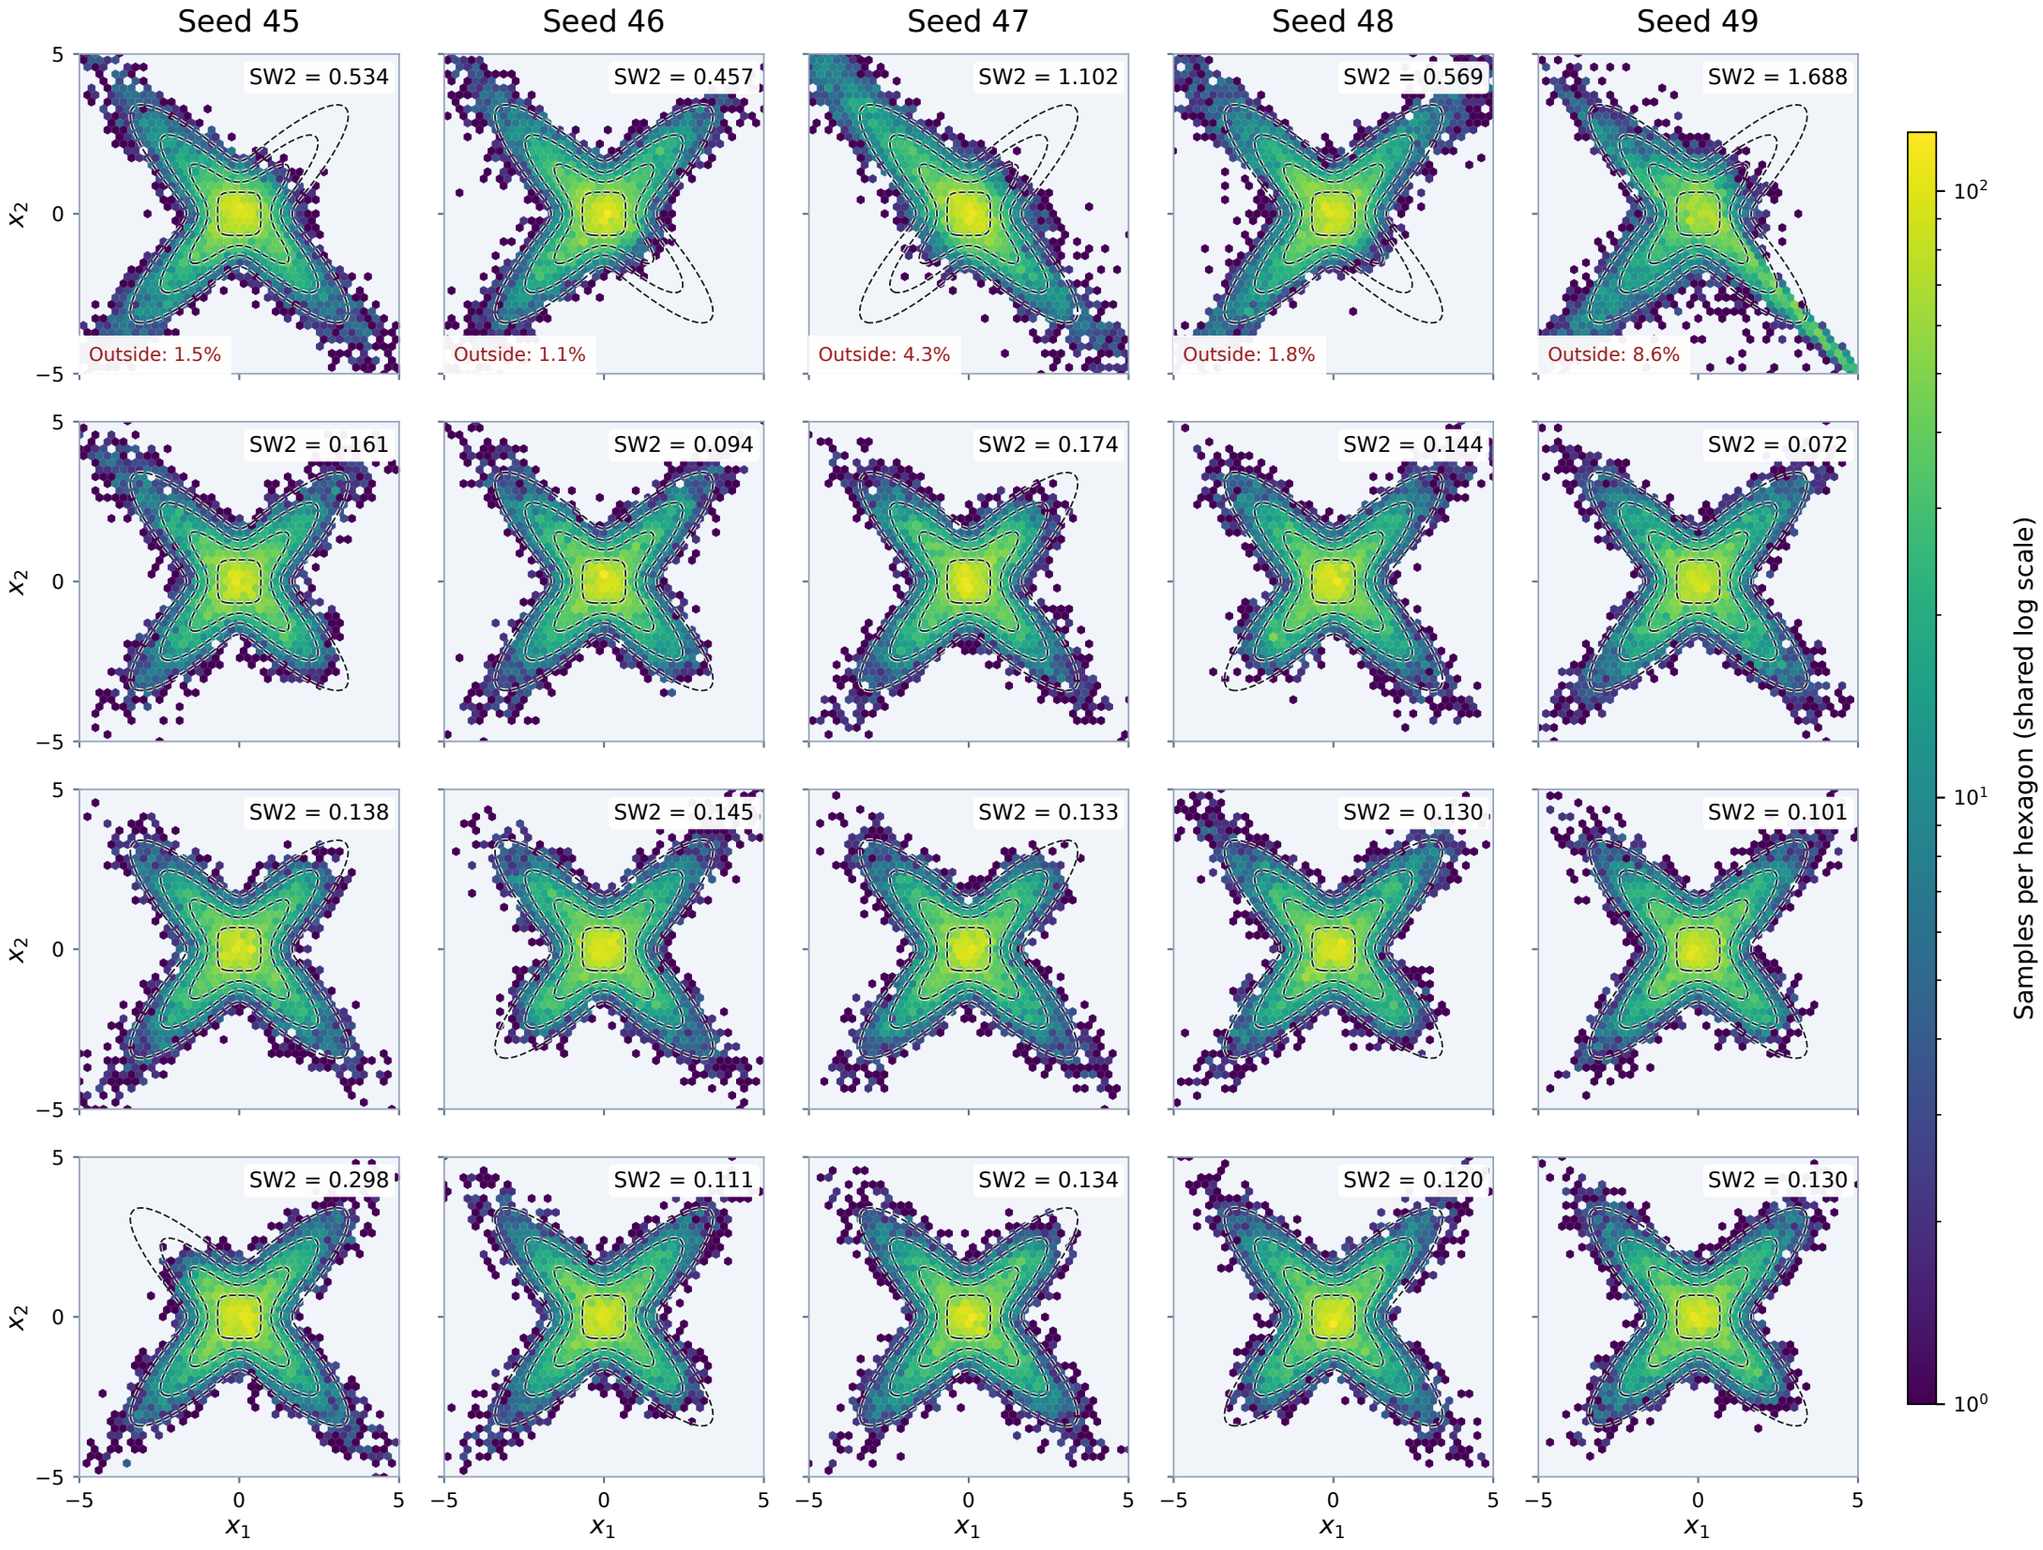
Mean SW2: 0.129 +/- 0.017

Global variance

One variance vector

$\eta_v = 10^{-3}$

Mean SW2: 0.159 +/- 0.079



10,000 saved samples per panel. Dashed contours show the exact X-shaped target.

Mean/backbone LR = 0.001; common schedule and axes. SW2 uses all samples; lower is better.

Row summaries: mean +/- sample SD across seeds. Off-panel sample fractions are noted if above 1%.